


Figure 14c · IR Instruction Count Is Not a Reliable Proxy for Dynamic Runtime

Panel A — IR count vs. runtime
(per-benchmark regression lines)



Panel B — Per-benchmark ρ (sorted)
25/30 have $|\rho| < 0.5$ | mean $\rho = -0.22$

