


IR Instruction Count is Not a Reliable Proxy for Dynamic Runtime

Panel A — IR count vs. runtime
(within-benchmark regression lines)



Panel B — Per-benchmark Spearman ρ
(25/30 benchmarks have $|\rho| < 0.5$)

